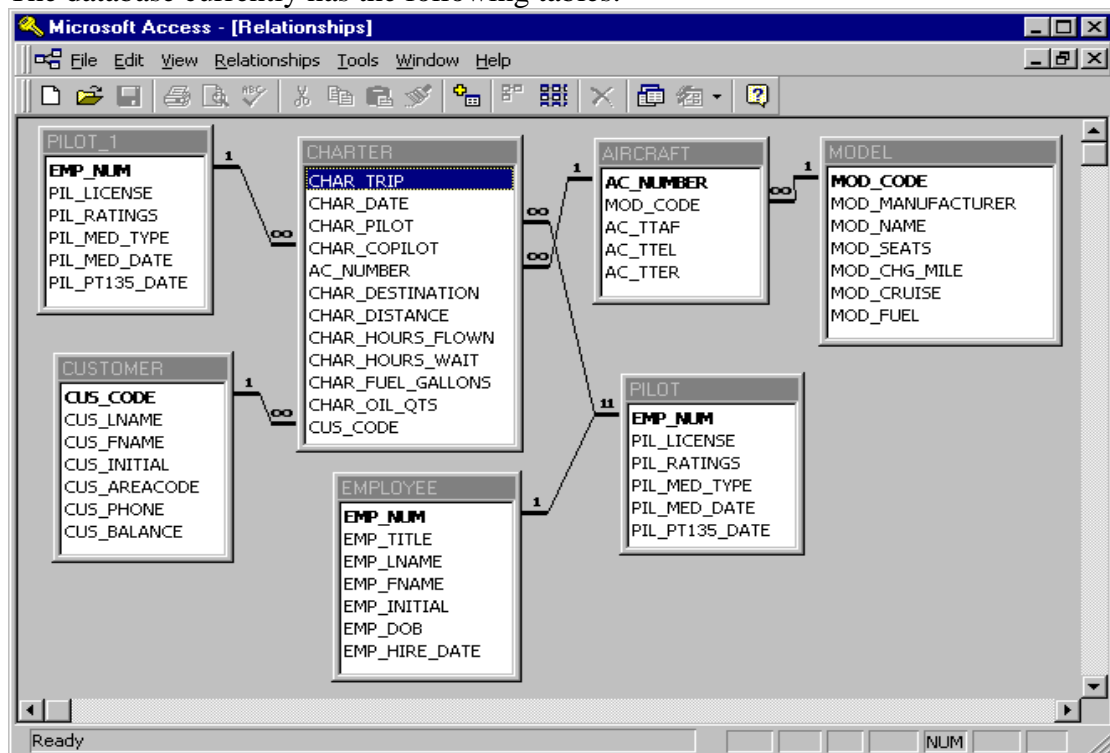


Data Exploration The Robcor Case Study

ROBCOR, Inc. provides “on demand” aviation charters, using a mix of different airplane and airplane types. Because ROBCOR, Inc., has grown rapidly, it has hired you to be its first database manager. Your first and critical assignment is to develop a decision support system to analyze the charter data. The charter operations manager wants to be able to analyze charter data such as total hours flown, total fuel used, and total revenue (charter distance x model charge per mile). She would also like to be able to drill-down by pilot, airplane model, and time periods. The main requirements for this database are to:

- Show the total revenue each month/year
- Show the total hours flown by each pilot
- Show the total fuel used by each airplane model.

The database currently has the following tables:



The tables can be copied from the dw account, using:

```
Select * from dw.<table_name>;
```

Or

```
Create Table <your_table_name> As  
Select ...  
From dw.<table_name>  
Where ...
```

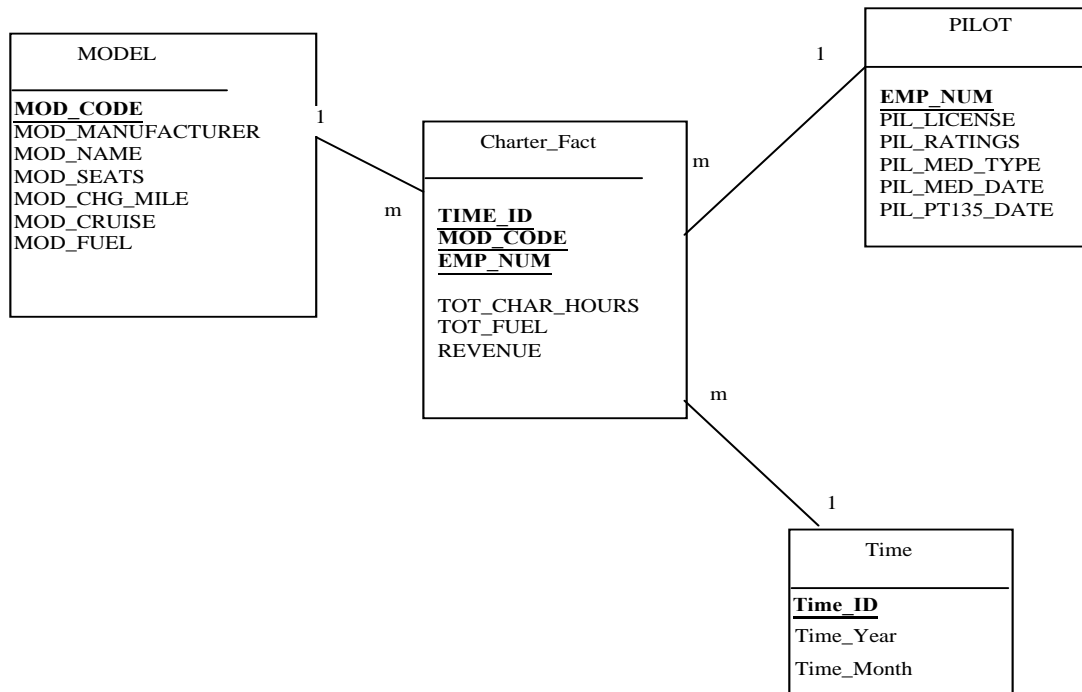
Given these requirements, complete the following:

- Create a star schema for the charter data.
- Define the dimensions and attributes for the charter operation’s star schema.
- Define the SQL statements for the implementation of the star schema.

4. Write the SQL statements to produce the following reports:
 - a. Show the total revenue each year
 - b. Show the total hours flown by each pilot
 - c. Show the total fuel used by each aircraft model

Solutions:

The Star Schema



--First create the dimensions

```

create table time As
select Distinct to_char(char_date, 'YYYYMM') as Time_ID,
               to_char(char_date, 'Month') as Time_Month,
               to_char(char_date, 'YYYY') as Time_Year
from dw.Charter;

```

```

create table model as
select * from dw.model;

```

```

create table pilot as
select * from dw.pilot;

```

--Second, create the Charter_fact (the fact table) table

```

create table charter_fact as
select C.Char_Pilot as EMP_Num,
       M.Mod_Code,
       to_char(C.Char_Date, 'YYYYMM') as Time_ID,
       sum(C.Char_Hours_Flown) as Tot_Char_Hours,
       sum(C.Char_Fuel_Gallons) as Tot_Fuel,
       sum(C.Char_Distance * M.Mod_chg_mile) as Revenue
from   dw.Charter C, dw.Model M, dw.Aircraft A
where  C.AC_Number=A.AC_Number and A.Mod_Code=M.Mod_Code
group by C.Char_Pilot, M.Mod_Code, to_char(C.Char_Date, 'YYYYMM');

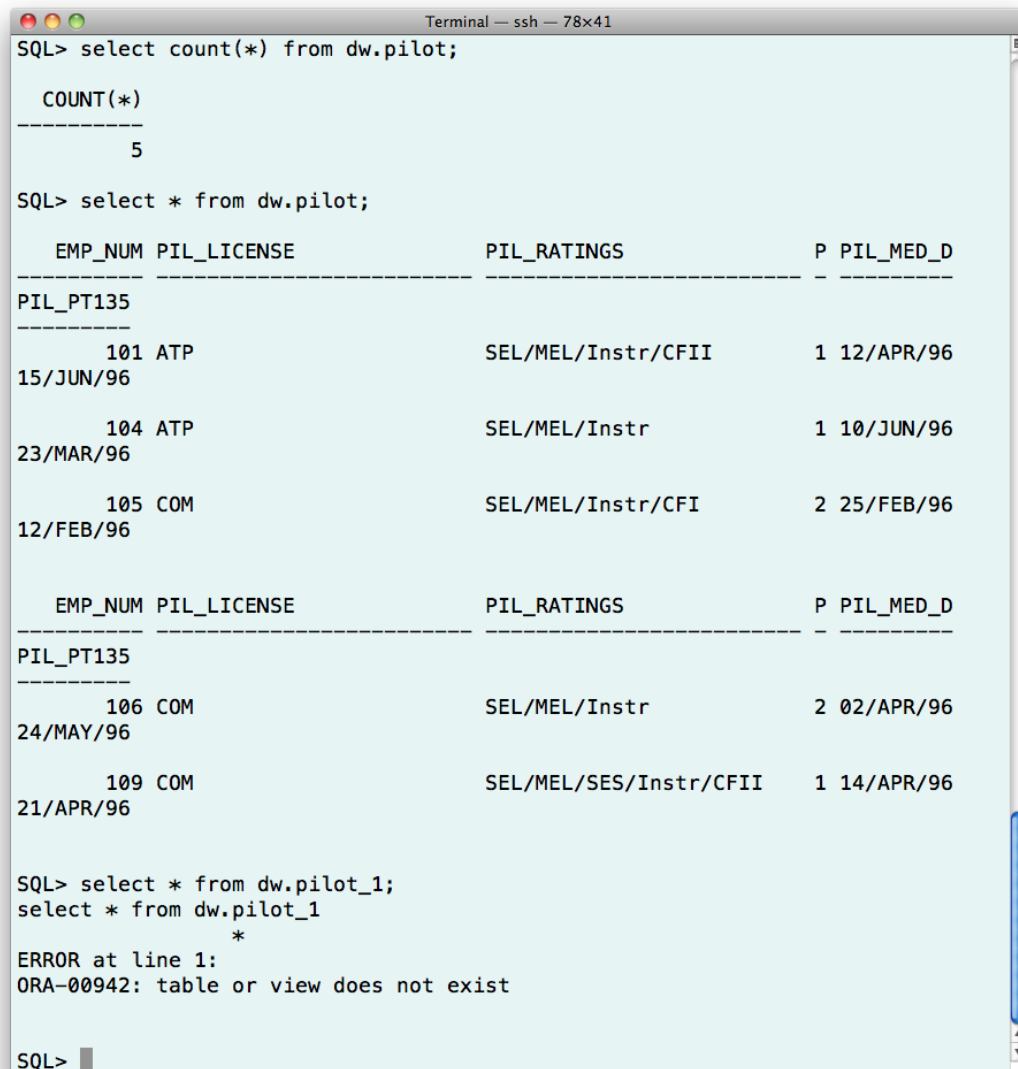
```

Now, we need to explore the data, both operational data, and data warehouse data. **Data exploration** is always necessary, because we will not be able to determine whether the results of our data warehouse are correct or not by just looking at the SQL above.

You first need to explore all tables in the operational database, including finding out what the operational data is about. Here is what I found:

1. Where is Pilot_1 table?

In the E/R diagram above, there is Pilot_1 entity. But I could not find the table. Pilot table is there, but not Pilot_1 table. There are five records in Pilot table.



```
Terminal — ssh — 78x41
SQL> select count(*) from dw.pilot;

  COUNT(*)
-----
         5

SQL> select * from dw.pilot;

  EMP_NUM  PIL_LICENSE  PIL_RATINGS  P  PIL_MED_D
-----
PIL_PT135
-----
      101  ATP          SEL/MEL/Instr/CFII  1  12/APR/96
15/JUN/96

      104  ATP          SEL/MEL/Instr      1  10/JUN/96
23/MAR/96

      105  COM          SEL/MEL/Instr/CFI   2  25/FEB/96
12/FEB/96

  EMP_NUM  PIL_LICENSE  PIL_RATINGS  P  PIL_MED_D
-----
PIL_PT135
-----
      106  COM          SEL/MEL/Instr      2  02/APR/96
24/MAY/96

      109  COM          SEL/MEL/SES/Instr/CFII  1  14/APR/96
21/APR/96

SQL> select * from dw.pilot_1;
select * from dw.pilot_1
                *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL>
```

However, when I do `Select * From dw.Charter`, the attributes are all complete, including Char_pilot and Char_copilot attributes.

```

Terminal — ssh — 73x19

SQL> desc dw.charter;

```

Name	Null?	Type
CHAR_TRIP		NUMBER(10)
CHAR_DATE		DATE
CHAR_PILOT		NUMBER(10)
CHAR_COPILOT		NUMBER(10)
AC_NUMBER		CHAR(5)
CHAR_DESTINATION		CHAR(3)
CHAR_DISTANCE		NUMBER(10)
CHAR_HOURS_FLOWN		FLOAT(126)
CHAR_HOURS_WAIT		FLOAT(126)
CHAR_FUEL_GALLONS		FLOAT(126)
CHAR_OIL_QTS		NUMBER(5)
CUS_CODE		NUMBER(10)

```

SQL>

```

After checking dw.charter table (select * from dw.charter), this is what I've got:

```

Terminal — ssh — 76x29

```

CHAR_TRIP	CHAR_DATE	CHAR_PILOT	CHAR_COPILOT	AC_NU	CHA	CHAR_DISTANCE
CHAR_HOURS_FLOWN	CHAR_HOURS_WAIT	CHAR_FUEL_GALLONS	CHAR_OIL_QTS	CUS_CODE		
10279	27/JUN/97	105		2289L	BNA	336
1.6		0		65.3		10010
10280	28/JUN/97	109		2289L	ATL	936
5.1		2.2		354.1	1	10011
10281	29/JUN/97	101		2778V	BNA	320
1.6		0		72.6	0	10016
10282	30/JUN/97	105		104 4278Y	GNV	1574
7.8		0		339.8	2	10014
10283	30/JUN/97	106		1484P	STL	472
2.9		4.9		97.2	1	10019

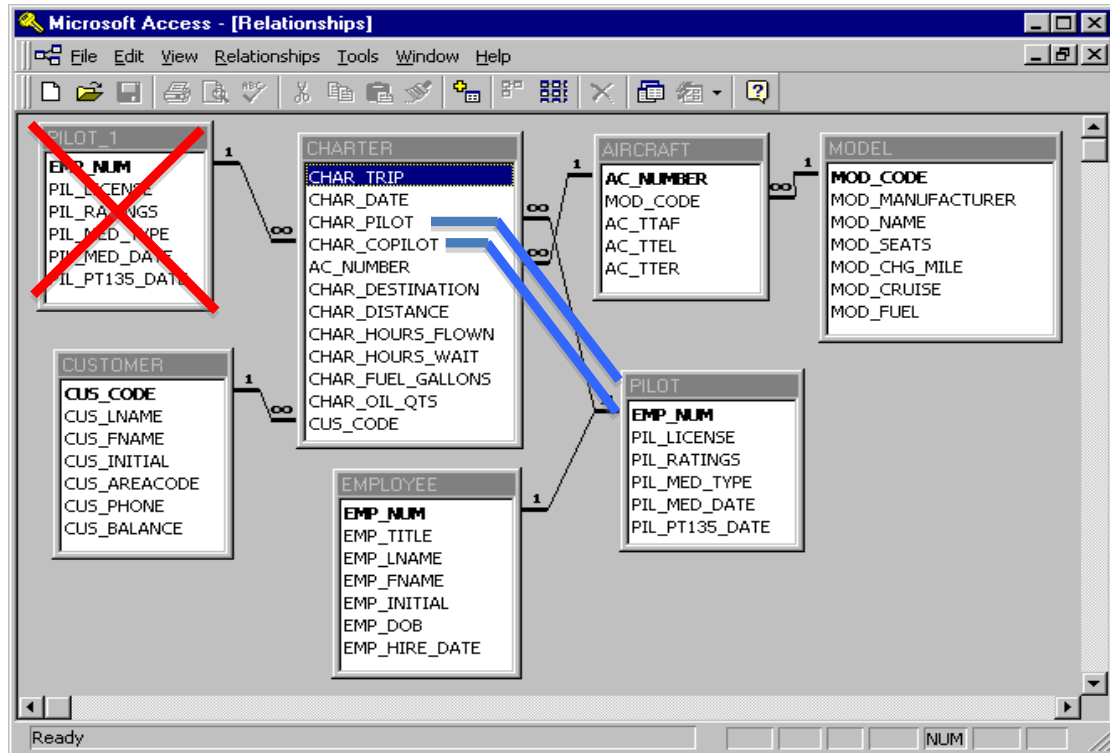
```

863 rows selected.

SQL>

```

It seems that Char_pilot and Char_copilot are drawing from the same table, that is table Pilot. Some records in dw.charter have entries in the Char_copilot column. So, the conclusion that I can draw is that the give E/R diagram is not accurate. It is perhaps the way the designer drew the E/R diagram was like that, to indicate that table Pilot is used twice by the Charter entity. Hence, a revised E/R diagram could be drawn like this:



While on the subject of pilot and copilot, let's explore more:

The first one is to check whether there are records where the copilot is the pilot (i.e. the entry in char_copilot is the same as in char_pilot). I found nothing. So this is correct. The second check is to find out whether there are records without pilot. Also nothing. So this is also correct.

```

Terminal — ssh — 81x24

SQL> select emp_num from dw.pilot;

  EMP_NUM
  -----
    101
    104
    105
    106
    109

SQL> select count(*) from dw.charter where char_pilot = char_copilot;

  COUNT(*)
  -----
         0

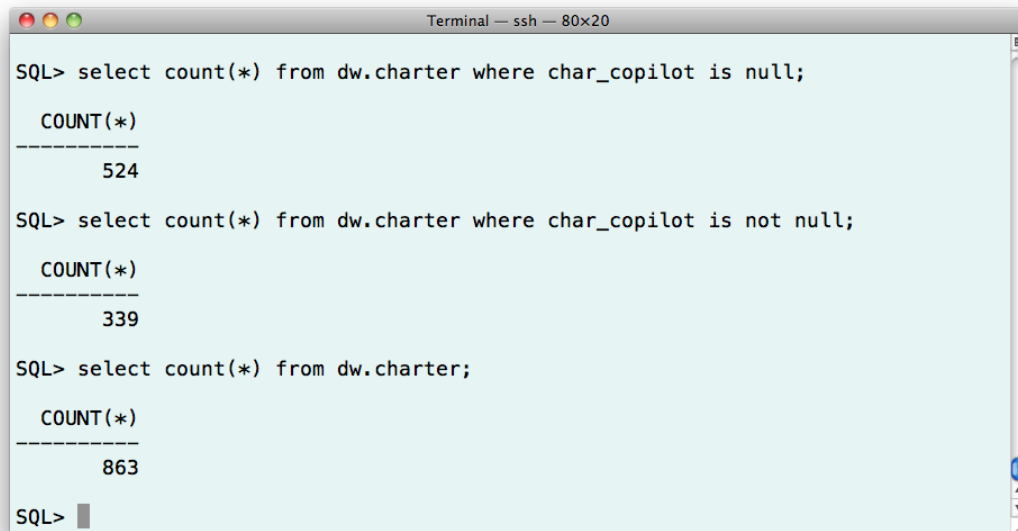
SQL> select count(*) from dw.charter where char_pilot is null;

  COUNT(*)
  -----
         0

SQL>

```

Now, let's check the copilots. We know that not all flights have copilots. Let's check how many flights with copilots, and how many flights without copilots, and see whether the sum of these match with the total number of records in dw.charter.



```
Terminal — ssh — 80x20

SQL> select count(*) from dw.charter where char_copilot is null;

  COUNT(*)
-----
       524

SQL> select count(*) from dw.charter where char_copilot is not null;

  COUNT(*)
-----
       339

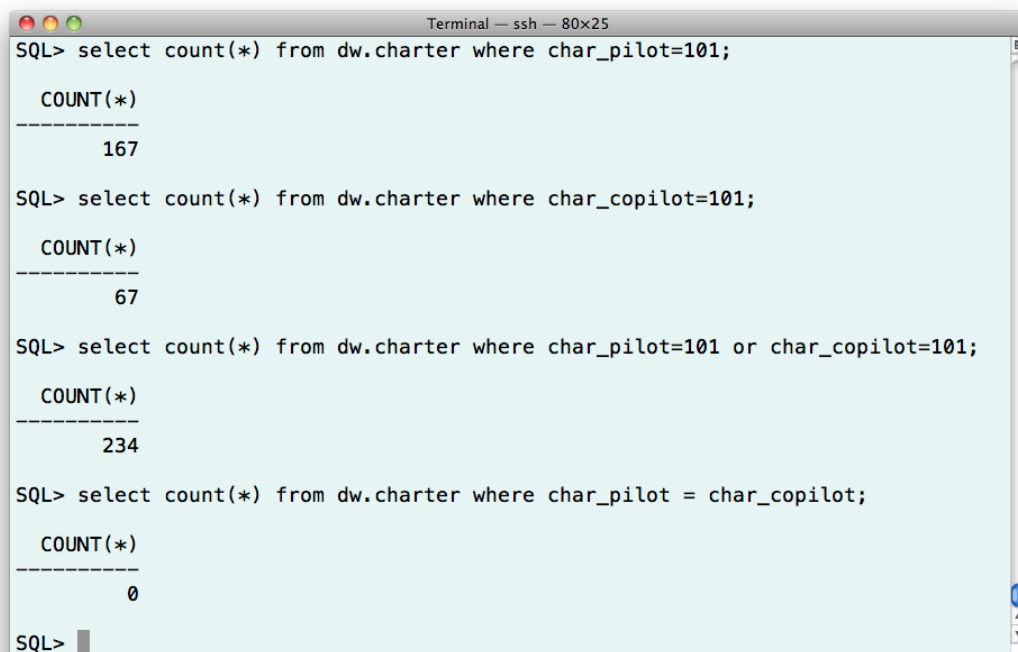
SQL> select count(*) from dw.charter;

  COUNT(*)
-----
       863

SQL> 
```

Everything about pilot and copilot seem to be fine. So, the only issue was with the E/R diagram.

The following shows number of charter flights for pilot 101 as a pilot, and as a copilot. The sum shows that, in this case, employee 101 cannot be pilot and a copilot for the same charter record. So, this is correct.



```
Terminal — ssh — 80x25

SQL> select count(*) from dw.charter where char_pilot=101;

  COUNT(*)
-----
       167

SQL> select count(*) from dw.charter where char_copilot=101;

  COUNT(*)
-----
        67

SQL> select count(*) from dw.charter where char_pilot=101 or char_copilot=101;

  COUNT(*)
-----
       234

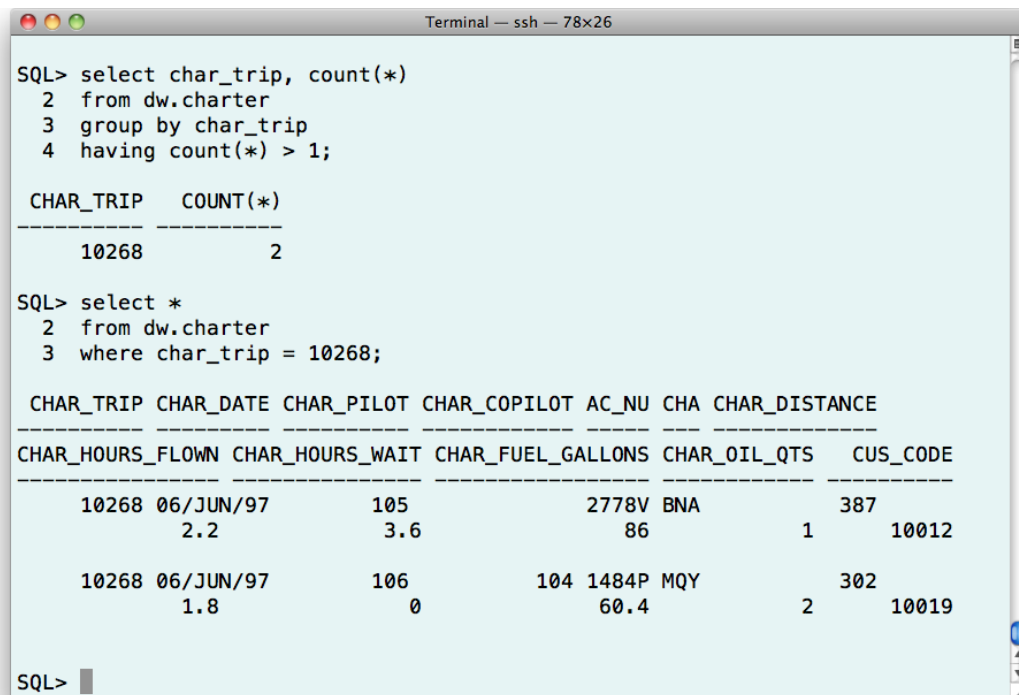
SQL> select count(*) from dw.charter where char_pilot = char_copilot;

  COUNT(*)
-----
         0

SQL> 
```

2. Duplicate Char_Trip

I found two different records in dw.charter sharing the same char_trip. By looking at these two records, it seems that these are two different charter records. Since in our data warehouse, char_trip does not affect the calculations in the fact table, we should not worry about this inconsistent data. So, the input data is not clean, but the unclean input data does not affect our data warehouse.



```
Terminal — ssh — 78x26

SQL> select char_trip, count(*)
2  from dw.charter
3  group by char_trip
4  having count(*) > 1;

CHAR_TRIP  COUNT(*)
-----
10268      2

SQL> select *
2  from dw.charter
3  where char_trip = 10268;

CHAR_TRIP  CHAR_DATE  CHAR_PILOT  CHAR_COPILOT  AC_NU  CHA  CHAR_DISTANCE
-----
10268  06/JUN/97      105      2778V BNA      387
2.2      3.6      86      1      10012

10268  06/JUN/97      106      104 1484P MQY      302
1.8      0      60.4      2      10019

SQL>
```

3. Should we consider copilot when calculating the fact table?

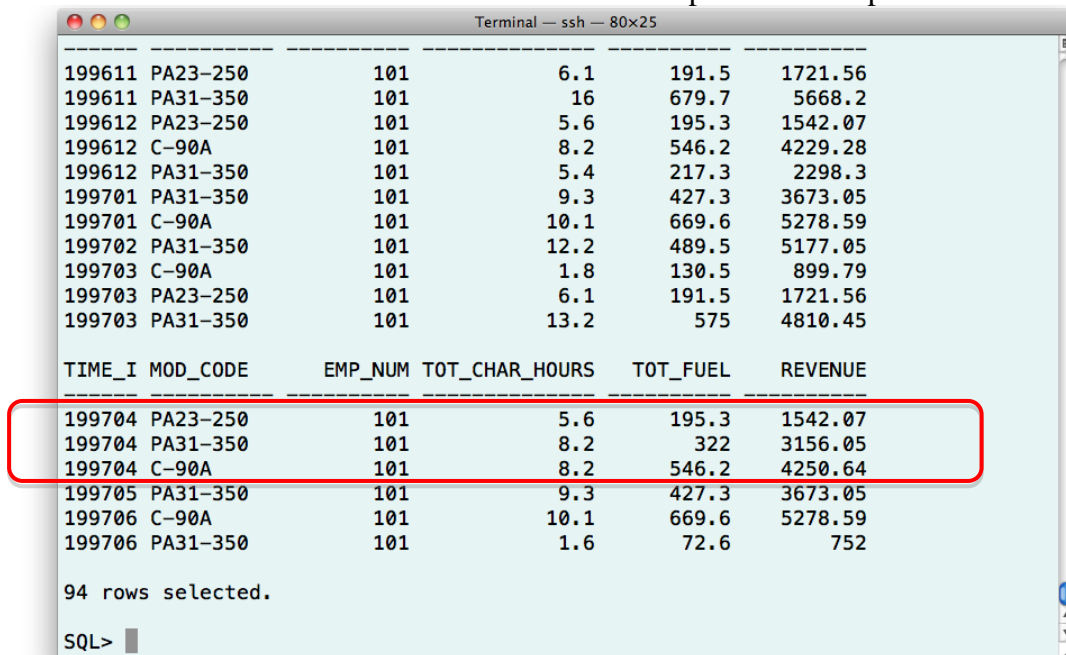
When Charter_fact is created, we do not consider copilot. The SQL to create Charter_fact is as follows:

```
create table charter_fact as
select C.Char_Pilot as EMP_Num,
       M.Mod_Code,
       to_char(C.Char_Date, 'YYYYMM') as Time_ID,
       sum(C.Char_Hours_Flown) as Tot_Char_Hours,
       sum(C.Char_Fuel_Gallons) as Tot_Fuel,
       sum(C.Char_Distance * M.Mod_chg_mile) as Revenue
from   dw.Charter C, dw.Model M, dw.Aircraft A
where  C.AC_Number=A.AC_Number and A.Mod_Code=M.Mod_Code
group by C.Char_Pilot, M.Mod_Code, to_char(C.Char_Date, 'YYYYMM');
```

Now let's explore pilot 101:

```
select *
from charter_fact
where emp_num=101
order by time_id;
```

The result is as follows. Look at the three records of pilot 101 in April 1997.



TIME_I	MOD_CODE	EMP_NUM	TOT_CHAR_HOURS	TOT_FUEL	REVENUE
199611	PA23-250	101	6.1	191.5	1721.56
199611	PA31-350	101	16	679.7	5668.2
199612	PA23-250	101	5.6	195.3	1542.07
199612	C-90A	101	8.2	546.2	4229.28
199612	PA31-350	101	5.4	217.3	2298.3
199701	PA31-350	101	9.3	427.3	3673.05
199701	C-90A	101	10.1	669.6	5278.59
199702	PA31-350	101	12.2	489.5	5177.05
199703	C-90A	101	1.8	130.5	899.79
199703	PA23-250	101	6.1	191.5	1721.56
199703	PA31-350	101	13.2	575	4810.45
199704	PA23-250	101	5.6	195.3	1542.07
199704	PA31-350	101	8.2	322	3156.05
199704	C-90A	101	8.2	546.2	4250.64
199705	PA31-350	101	9.3	427.3	3673.05
199706	C-90A	101	10.1	669.6	5278.59
199706	PA31-350	101	1.6	72.6	752

94 rows selected.

SQL>

Now let's check employee 101 as pilot and copilot in April 1997:

```
select *
from dw.charter
where to_char(char_date, 'YYYYMM') = '199704'
and (char_pilot = 101 or char_copilot = 101);
```


CHAR_TRIP	CHAR_DATE	CHAR_PILOT	CHAR_COPILOT	AC_NU	CHA	CHAR_DISTANCE	
CHAR_HOURS_FLOWN	CHAR_HOURS_WAIT	CHAR_FUEL_GALLONS	CHAR_OIL_QTS	CUS_CODE			
10229	03/APR/97	101	109	4278Y	BNA	365	
2.8		0	104.7		0	10010	
10233	12/APR/97	101	109	1484P	MOB	799	
5.6		0	195.3		0	10018	
10235	14/APR/97	105	101	2778V	ATL	982	
5.6		5.4	241.1		1	10017	
CHAR_TRIP	CHAR_DATE	CHAR_PILOT	CHAR_COPILOT	AC_NU	CHA	CHAR_DISTANCE	
CHAR_HOURS_FLOWN	CHAR_HOURS_WAIT	CHAR_FUEL_GALLONS	CHAR_OIL_QTS	CUS_CODE			
10239	19/APR/97	101		2289L	GNV	1592	
8.2		0		546.2	0	10014	
10245	28/APR/97	101		4278Y	ATL	978	
5.4		2.8		217.3	0	10016	

The result shows that employee 101 was a copilot in one charter flight in April 1997.

The fact table for employee 101 does not include the data of employee 101 as a copilot. You can double-check the attributes of tot_char_hours, tot_fuel_used, and revenue in the fact table. One way to solve this problem is to create another fact table solely for copilot. Hence we need Charter_fact2:

```
create table charter_fact2 as
select C.Char_coPilot as EMP_Num,
       M.Mod_Code,
       to_char(C.Char_Date, 'YYYYMM') as Time_ID,
       sum(C.Char_Hours_Flown) as Tot_Char_Hours,
       sum(C.Char_Fuel_Gallons) as Tot_Fuel,
       sum(C.Char_Distance * M.Mod_chg_mile) as Revenue
from   dw.Charter C, dw.Model M, dw.Aircraft A
where  C.AC_Number=A.AC_Number and A.Mod_Code=M.Mod_Code
group by C.Char_Pilot, M.Mod_Code, to_char(C.Char_Date, 'YYYYMM');
```

Let's see the records of copilot 101:

```
select *
from charter_fact2
where emp_num=101
order by time_id;
```

```
Terminal — ssh — 75x20
199608 C-90A          101          6.7          459.5         4392.15
199609 PA31-350       101          12.6          516.3         5428.5
199610 PA23-250       101           4.1          131.6         1269.94
199611 PA31-350       101           5.6          223.9         2354.7
199612 PA31-350       101           5.6          241.1         2307.7
199612 PA23-250       101           4.1          131.6         1269.94
199701 PA31-350       101          12.6          516.3         5428.5
199701 PA23-250       101           3.5           124          924.47
```

TIME_I	MOD_CODE	EMP_NUM	TOT_CHAR_HOURS	TOT_FUEL	REVENUE
199702	PA23-250	101	4.1	131.6	1269.94
199703	PA31-350	101	13.5	531.4	5412.05
199704	PA31-350	101	5.6	241.1	2307.7
199705	PA31-350	101	12.6	516.3	5428.5
199705	PA23-250	101	7.6	255.6	2194.41

49 rows selected.

SQL>

So, now we have two star schemas: one for pilot, and the other for co-pilot. The dimensions and fact for these two star schemas are the same, but the contents of the fact table are different. Let's see first how many records are there in both fact tables.

```

Terminal — ssh — 77x26

SQL> select count(*) from charter_fact;

  COUNT(*)
-----
      480

SQL> select count(*) from charter_fact2;

  COUNT(*)
-----
      390

SQL> select count(*) from charter_fact where emp_num=101;

  COUNT(*)
-----
       94

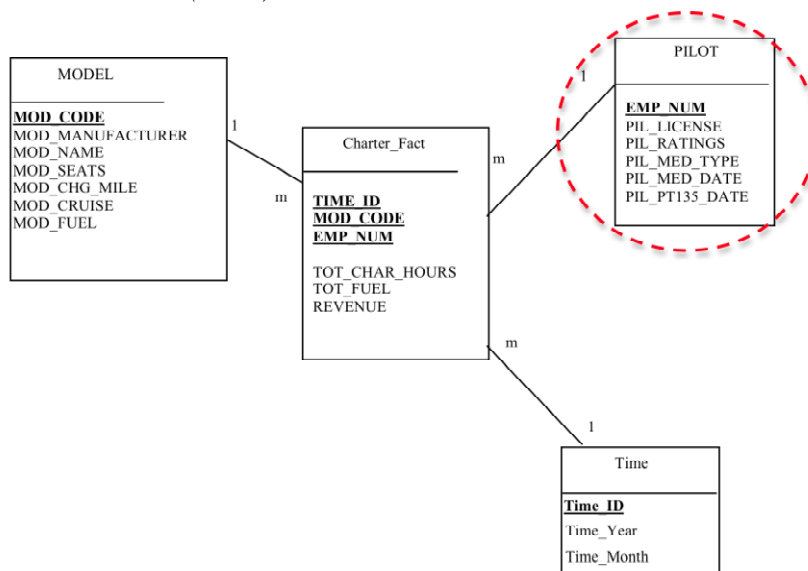
SQL> select count(*) from charter_fact2 where emp_num=101;

  COUNT(*)
-----
       49

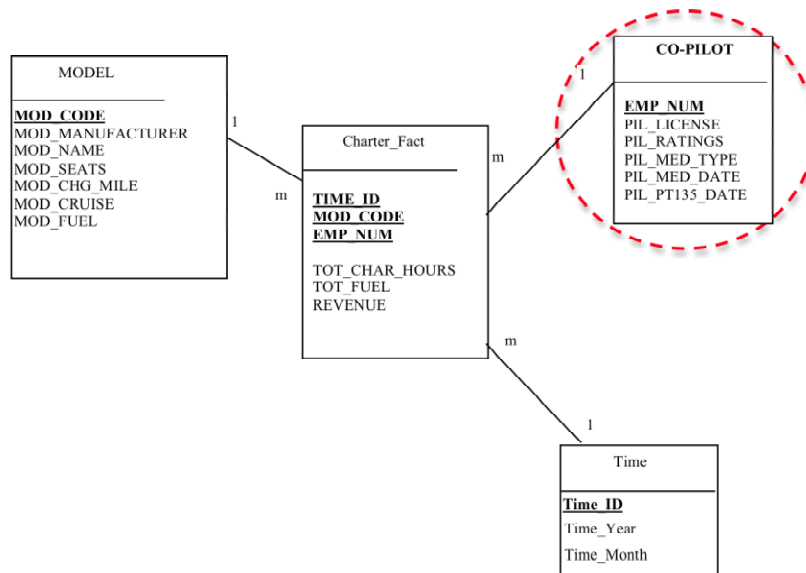
SQL>

```

Star Schema-1 (Pilot)



Star Schema-2 (Co-Pilot):



4. Is it possible to combine the two star schemas (the two facts)?

It is now possible to have a third fact table, which combines both pilot and co-pilot fact measures. So, what we would like to do is to “merge” the two charter facts, so that the new fact table incorporates both pilots and copilots.

To merge the two fact tables, we can use a UNION command in SQL, something like:

```

select * from charter_fact
union
select * from charter_fact2
  
```

We can put this UNION command in a CREATE TABLE AS SELECT command as follows:

```

create table charter_fact3 as
select time_id,
       mod_code,
       emp_num,
       sum(tot_char_hours) as tot_char_hours,
       sum(tot_fuel) as tot_fuel,
       sum(revenue) as revenue
from (
  select * from charter_fact
  union
  select * from charter_fact2)
group by time_id, mod_code, emp_num;
  
```

It basically merging the two fact tables, and then the result of this merging or union process is grouped by time_id, mod_code, and emp_num, so that records having the same time_id, mod_code, and emp_num are grouped into one record, where the new sum is then calculated.

The following shows number of records in charter_fact, charter_fact2, and charter_fact3 tables:

```
Terminal — ssh — 75x20

SQL> select count(*) from charter_fact;

  COUNT(*)
-----
      480

SQL> select count(*) from charter_fact2;

  COUNT(*)
-----
      390

SQL> select count(*) from charter_fact3;

  COUNT(*)
-----
      686

SQL>
```

Now, let's check employee 101 as a case study:

```
select *
from charter_fact
where emp_num=101
order by time_id;
```

```
Terminal — ssh — 73x15

199703 PA23-250      101      6.1      191.5      1721.56
199703 PA31-350      101     13.2       575      4810.45

TIME_I MOD_CODE      EMP_NUM TOT_CHAR_HOURS  TOT_FUEL  REVENUE
-----
199704 PA23-250      101       5.6      195.3      1542.07
199704 PA31-350      101       8.2       322      3156.05
199704 C-90A         101       8.2      546.2      4250.64
199705 PA31-350      101       9.3      427.3      3673.05
199706 C-90A         101      10.1      669.6      5278.59
199706 PA31-350      101       1.6       72.6        752

94 rows selected.

SQL>
```

And now employee 101 in charter_fact2 (as a copilot):

```
select *
from charter_fact2
where emp_num=101
order by time_id;
```

```

Terminal — ssh — 73x13
199701 PA23-250      101      3.5      124      924.47
TIME_I MOD_CODE      EMP_NUM TOT_CHAR_HOURS  TOT_FUEL  REVENUE
-----
199702 PA23-250      101      4.1      131.6    1269.94
199703 PA31-350      101     13.5     531.4    5412.05
199704 PA31-350      101      5.6     241.1    2307.7
199705 PA31-350      101     12.6     516.3    5428.5
199705 PA23-250      101      7.6     255.6    2194.41

49 rows selected.

SQL>

```

And finally, employee 101 in charter_fact3:

```

select *
from charter_fact3
where emp_num=101
order by time_id;

```

```

Terminal — ssh — 73x16
199703 PA23-250      101      6.1     191.5    1721.56
TIME_I MOD_CODE      EMP_NUM TOT_CHAR_HOURS  TOT_FUEL  REVENUE
-----
199703 PA31-350      101     26.7    1106.4   10222.5
199704 PA23-250      101      5.6     195.3    1542.07
199704 PA31-350      101     13.8     563.1    5463.75
199704 C-90A         101      8.2     546.2    4250.64
199705 PA23-250      101      7.6     255.6    2194.41
199705 PA31-350      101     21.9     943.6    9101.55
199706 C-90A         101     10.1     669.6    5278.59
199706 PA31-350      101      1.6      72.6      752

107 rows selected.

SQL>

```

Notice that the calculation of tot_char_hours, tot_fuel, and revenue now incorporates both pilot and copilot data.

Is this correct?

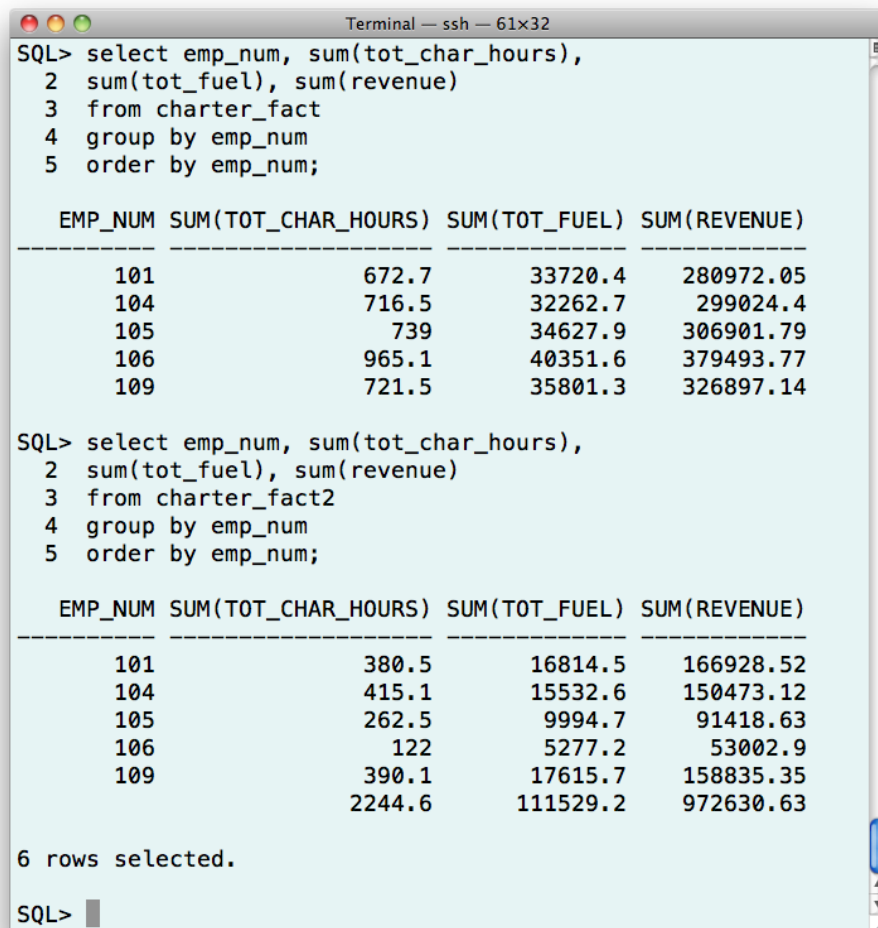
The answer is No!

Tot_char_hourss in charter_fact3, as shown in the above screenshot, is correct only in the context of Emp_num, which in this case it is pilot and co-pilot. In the above example, Emp_num 101, is a pilot in April 1997 flying aircraft model PA31-350 accumulated 8.2 hours, and the same pilot on the same month flying the same aircraft model as a copilot accumulated 5.6 hours, with a total of 13.8 hours. So, this is correct.

However, if you look at from the month point of view (e.g. April 1997), total hours of 13.8 is a double-dip, because when Emp_num 101 is a co-pilot, the flying hour is also recorded by the pilot of that flight (which is not Emp_num 101). Therefore, it is not possible to just merge the two facts.

However, we are able to merge to two facts, if we only take Pilot/Co-Pilot dimension, and not the other two dimensions.

The following are the fact measures of pilots and co-pilots from the two facts.



```

Terminal — ssh — 61x32
SQL> select emp_num, sum(tot_char_hours),
2  sum(tot_fuel), sum(revenue)
3  from charter_fact
4  group by emp_num
5  order by emp_num;

  EMP_NUM  SUM(TOT_CHAR_HOURS)  SUM(TOT_FUEL)  SUM(REVENUE)
-----
      101             672.7        33720.4      280972.05
      104             716.5        32262.7      299024.4
      105              739        34627.9      306901.79
      106            965.1        40351.6      379493.77
      109            721.5        35801.3      326897.14

SQL> select emp_num, sum(tot_char_hours),
2  sum(tot_fuel), sum(revenue)
3  from charter_fact2
4  group by emp_num
5  order by emp_num;

  EMP_NUM  SUM(TOT_CHAR_HOURS)  SUM(TOT_FUEL)  SUM(REVENUE)
-----
      101             380.5        16814.5      166928.52
      104             415.1        15532.6      150473.12
      105             262.5         9994.7       91418.63
      106              122         5277.2        53002.9
      109             390.1        17615.7      158835.35
                2244.6        111529.2      972630.63

6 rows selected.

SQL>

```

Now, if we merge the above two facts, but only taking emp_num dimension, then the result is as follows:

```

create table charter_fact3b as
select emp_num,
       sum(tot_char_hours) as tot_char_hours,
       sum(tot_fuel) as tot_fuel,
       sum(revenue) as revenue
from (
  select * from charter_fact
  union
  select * from charter_fact2)
group by emp_num
order by emp_num;

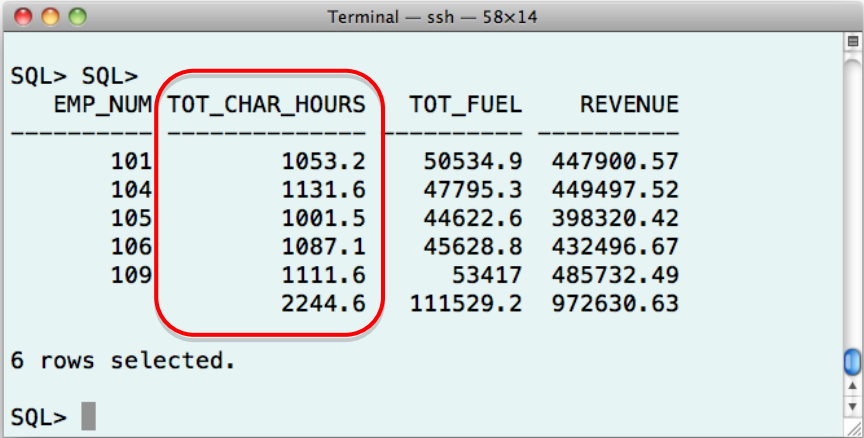
```

Or you could do this:

```
create table charter_fact3b as
select emp_num,
       sum(tot_char_hours) as tot_char_hours,
       sum(tot_fuel) as tot_fuel,
       sum(revenue) as revenue
from (
  select emp_num,
         sum(tot_char_hours) as tot_char_hours,
         sum(tot_fuel) as tot_fuel,
         sum(revenue) as revenue
  from charter_fact
  group by emp_num
 union
  select emp_num,
         sum(tot_char_hours) as tot_char_hours,
         sum(tot_fuel) as tot_fuel,
         sum(revenue) as revenue
  from charter_fact2
  group by emp_num)
group by emp_num
order by emp_num;
```

What is the difference between the two create table charter_fact3b methods?
The result will still be the same.

```
select * from charter_fact3b;
```



SQL> SQL>	EMP_NUM	TOT_CHAR_HOURS	TOT_FUEL	REVENUE
	101	1053.2	50534.9	447900.57
	104	1131.6	47795.3	449497.52
	105	1001.5	44622.6	398320.42
	106	1087.1	45628.8	432496.67
	109	1111.6	53417	485732.49
		2244.6	111529.2	972630.63

6 rows selected.

SQL>

If we look at Total Charter Hours in Fact3, this column shows the total hours for each employee, whether the employee is a pilot or a co-pilot. For example, for Employee 101, the total charter hours is 1053.2 hours (as a pilot or as a co-pilot). This is because we take the data from charter_fact (for pilot), union it with charter_fact2 (for co-pilot), and re-calculate the total charter hours.

However, if we look at the other two fact measures: Total Fuel and Total Revenue, the analysis will be misleading. For example, it is true that Employee 101 (as a pilot

or as a co-pilot) used up 50534.9 gallons of fuel. However, some of the trips that 101 made as a pilot, had someone else as a co-pilot (e.g. Employee 105). But Employee 105 also calculated the trip that he made as a co-pilot in the total fuel calculation; hence, the calculation of the total fuel is doubled; and is therefore incorrect.

The same mistake also happens for the Total Revenue. It is try that Employee 101 (as a pilot and as a co-pilot) brings a total revenue of \$447900.57. However, some of the trips that Employee 101 made are together with another employee. As a result, the revenue for that particular trip will be counted twice; and is therefore incorrect.

So, Total Fuel Used and Total Revenue fact measures should be excluded from charter_fact3. This raises two questions:

1. Why is total charter hours correct, whereas total fuel and total revenue are incorrect?
2. Is it still necessary to create charter_fact3 with employee as the only dimension and total charter hours as the only fact measure?

The answer to question 1 is that total charter hours are **individual**. Each employee has his/her own flying hours, and even when they fly together as a pilot and co-pilot, the flying hours are still individual.

On the other hand, total fuel used and total revenue are **not individual** – they are related to the company. So for example, in one trip, that has a pilot and a co-pilot, has \$1000 of revenue and use 100 gallons of fuel, it does not mean that the revenue of the pilot is \$1000 and the revenue of the co-pilot is another \$1000; a total of \$2000. The revenue is still \$1000, regardless whether there is a pilot only on the trip, or a pilot and a co-pilot on that trip. The same principal is applied to the fuel used.

In order to answer question 2, we need to evaluate how to create charter_fact3 correct. This is the charter_fact3 having emp_num and tot_char_hours only.

```
create table charter_fact3c as
select emp_num, sum(tot_char_hours) as tot_char_hours
from (
    select * from charter_fact
    union
    select * from charter_fact2)
group by emp_num
order by emp_num;
```

However, the use of charter_fact3 is rather limited. We could query the first two charter facts directly, using the following SQL:

```
select emp_num, sum(tot_char_hours) as tot_char_hours
from (
    select * from charter_fact
    union
    select * from charter_fact2)
group by emp_num
order by emp_num;
```

Therefore, the first two facts are sufficient: one star schema for pilot, and another star schema for co-pilot.

The main reason why it is not necessary to create a combined fact, is because the transaction (e.g. a charter trip), can have two “personnel”. Therefore, we need to have two star schema to capture the roles of these two personnel: pilot and co-pilot. Combining two star schemas into one is generally not possible or with a limited use.

On the other hand, if the transaction records have one personnel or pilot (for example, there is no notion of co-pilot), but there are two types of pilot: permanent pilot and sessional pilot; if we have two star schemas: one for permanent pilot and another for sessional pilot, these two star schemas can be combined by using a simple union operator, without any re-calculation of the fact measures. Because, for example, the total revenue of a permanent pilot will not overlap with the total revenue of a sessional pilot. Combining these two records in a new fact will remain two records.